

Activity: Visualizing Ridge Regression

Tutorial-2

Objective

Understand how **Ridge Regression (L2 Regularization)** prevents overfitting by visualizing how the penalty term α (alpha) "tames" the model's complexity.

Part 1: The Setup

We will fit a high-degree polynomial (Degree 10) to a noisy sine wave. A standard Linear model will try to hit every noise point, creating wild "wiggles." We will use Ridge Regression to smooth it out.

Instructions: Copy and run the following code in your Python environment (Jupyter Notebook or Colab).

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.pipeline import make_pipeline
4 from sklearn.preprocessing import PolynomialFeatures
5 from sklearn.linear_model import Ridge, LinearRegression
6
7 # 1. Generate Synthetic Data (A sine wave with noise)
8 np.random.seed(42)
9 X = np.sort(np.random.rand(20) * 5)[:, np.newaxis] # 20 random points
10 y = np.sin(X).ravel() + np.random.normal(0, 0.2, 20) # True function + noise
11
12 # 2. Define a Test Set (Smooth line for plotting)
13 X_plot = np.linspace(0, 5, 100)[:, np.newaxis]
14
15 # 3. Compare 3 Models
16 plt.figure(figsize=(10, 6))
17 plt.scatter(X, y, color='black', label='Noisy Data')
18
19 # Model A: Standard Linear Regression (No Penalty)
20 # We use degree=10 to force the model to overfit
21 model_ols = make_pipeline(PolynomialFeatures(10), LinearRegression())
22 model_ols.fit(X, y)
23 plt.plot(X_plot, model_ols.predict(X_plot), color='red', linewidth=2, label='OLS (No Penalty)')
24
25 # Model B: Ridge with Low Penalty
26 model_ridge_low = make_pipeline(PolynomialFeatures(10), Ridge(alpha=0.1))
27 model_ridge_low.fit(X, y)
28 plt.plot(X_plot, model_ridge_low.predict(X_plot), color='blue', linestyle='--', label='Ridge (alpha=0.1)')
29
30 # Model C: Ridge with High Penalty
31 model_ridge_high = make_pipeline(PolynomialFeatures(10), Ridge(alpha=10.0))
32 model_ridge_high.fit(X, y)
33 plt.plot(X_plot, model_ridge_high.predict(X_plot), color='green', label='Ridge (alpha=10)')
34
```

```

35 plt.title('Taming the Wiggles: OLS vs Ridge')
36 plt.ylim(-2, 2)
37 plt.legend()
38 plt.show()

```

Part 2: Exploration Questions

Run the code above and analyze the graph to answer the following:

1. **The "Explosion" (Variance):** Look at the **Red Line** (OLS).
 - Why does it curve wildly at the edges compared to the Green line?
 - Does this model have High Bias or High Variance?
2. **The "Flatline" (Bias):**
 - Modify the code for **Model C**. Change `alpha=10.0` to `alpha=1000.0`.
 - Re-run the code. What happens to the Green line? Why?
3. **The "Sweet Spot":**
 - Try to find an `alpha` value for **Model B** (Blue dashed line) that fits the pattern smoothly without being too flat or too wiggly.
 - Optimal Alpha found: _____

Part 3:

Why did the Red line wiggle so much? It used massive coefficients to twist the line. Run this snippet to see the difference in the "budget" used by the models:

```

1 print("Sum of squared coefficients (OLS):", (model_ols.steps[1][1].coef_**2) .
      sum())
2 print("Sum of squared coefficients (Ridge):", (model_ridge_high.steps[1][1] .
      coef_**2).sum())

```

- **OLS Budget:** _____ (Likely Huge)
- **Ridge Budget:** _____ (Likely Small)